

```
In [147]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn import preprocessing
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, roc_auc_score
from imblearn.over_sampling import SMOTE
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

Data PreprPreprocessing

```
In [48]: disease = pd.read_csv("framingham.csv")
```

```
In [49]: disease
```

```
Out[49]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp
0	1	39	4.0	0	0.0	0.0	0	C
1	0	46	2.0	0	0.0	0.0	0	C
2	1	48	1.0	1	20.0	0.0	0	C
3	0	61	3.0	1	30.0	0.0	0	1
4	0	46	3.0	1	23.0	0.0	0	C
...
4235	0	48	2.0	1	20.0	NaN	0	C
4236	0	44	1.0	1	15.0	0.0	0	C
4237	0	52	2.0	0	0.0	0.0	0	C
4238	1	40	3.0	0	0.0	0.0	0	1
4239	0	39	3.0	1	30.0	0.0	0	C

4240 rows × 16 columns



Renaming the male column to Sex_male to know that 1 is male and 0 is female


```
In [50]: disease.rename(columns={'male' : 'Sex_male'}, inplace = True)
```

In [51]: `disease`

Out[51]:

	Sex_male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevale
0	1	39	4.0	0	0.0	0.0	0	
1	0	46	2.0	0	0.0	0.0	0	
2	1	48	1.0	1	20.0	0.0	0	
3	0	61	3.0	1	30.0	0.0	0	
4	0	46	3.0	1	23.0	0.0	0	
...
4235	0	48	2.0	1	20.0	NaN	0	
4236	0	44	1.0	1	15.0	0.0	0	
4237	0	52	2.0	0	0.0	0.0	0	
4238	1	40	3.0	0	0.0	0.0	0	
4239	0	39	3.0	1	30.0	0.0	0	

4240 rows × 16 columns



Handling the missing values

In [54]: `disease.isnull().sum()`

Out[54]:

Sex_male	0
age	0
currentSmoker	0
cigsPerDay	29
BPMeds	53
prevalentStroke	0
prevalentHyp	0
diabetes	0
totChol	50
sysBP	0
diaBP	0
BMI	19
heartRate	1
glucose	388
TenYearCHD	0
dtype:	int64

In [55]: `disease['cigsPerDay'].skew()`

Out[55]: 1.2470523561848126

```
In [56]: #As the skew of cigsPerDay is greater than 0 so we fill the null value of t
disease['cigsPerDay'] = disease['cigsPerDay'].fillna(disease['cigsPerDay'].

Out[57]: 0.8718805634765354

In [58]: #As the skew is 0.87 which is greater than 0 so we fill the null value of t
disease['totChol']=disease['totChol'].fillna(disease['totChol'].mean())

In [59]: #As the BPMeds is in binary form so we use mode to fill the null values
disease['BPMeds']=disease['BPMeds'].fillna(disease['BPMeds'].mode()[0])

In [60]: disease['BMI'].skew()

Out[60]: 0.9821832986950597

In [61]: #As the skew is 0.98 which is greater than 0 so we fill the null value of E
disease['BMI']=disease['BMI'].fillna(disease['BMI'].median())

In [62]: disease['glucose'].skew()

Out[62]: 6.2149483495346765

In [63]: #As the skew is 6.214 which is highly greater than 0 so we fill the null va
disease['glucose']=disease['glucose'].fillna(disease['glucose'].median())

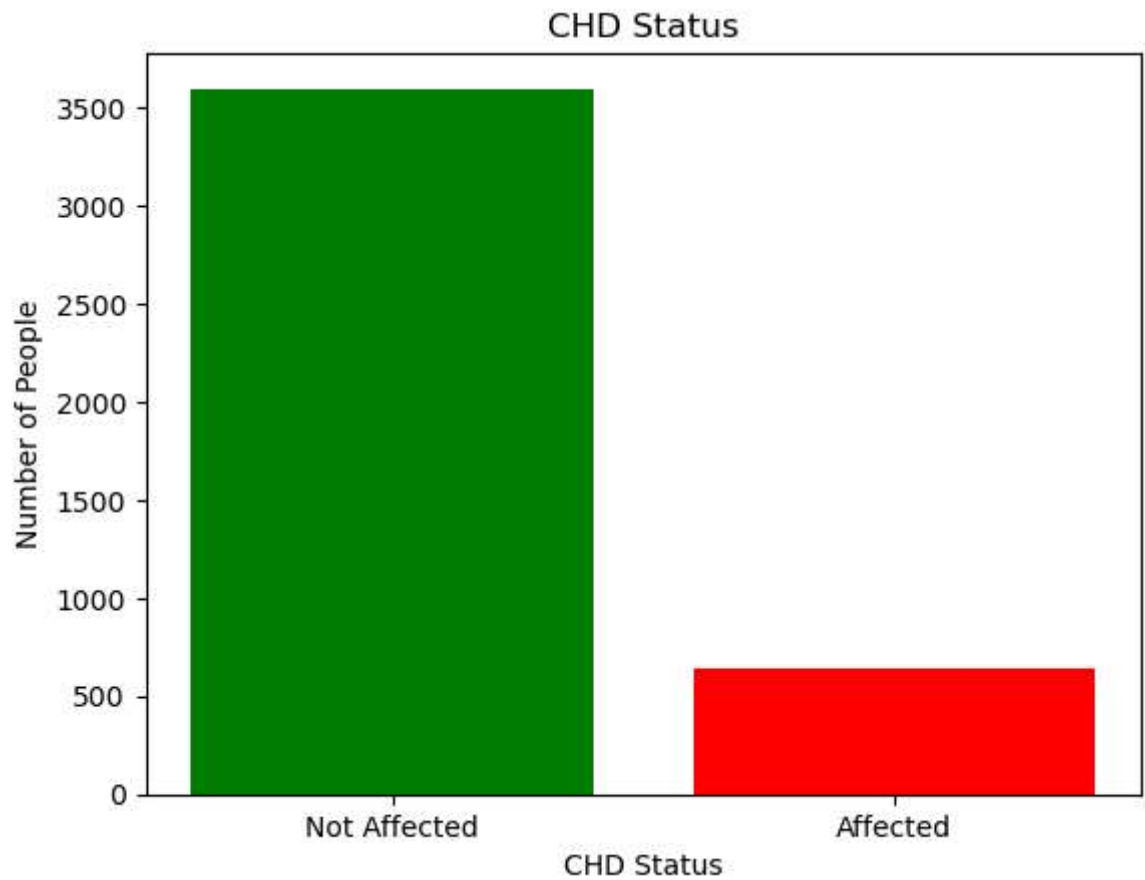
In [64]: disease.isnull().sum()

Out[64]: Sex_male          0
age          0
currentSmoker  0
cigsPerDay    0
BPMeds        0
prevalentStroke  0
prevalentHyp   0
diabetes       0
totChol        0
sysBP          0
diaBP          0
BMI            0
heartRate      1
glucose        0
TenYearCHD     0
dtype: int64
```

Making a Bar Plot to Count people affected by CHD

```
In [65]: counts = disease["TenYearCHD"].value_counts()

plt.bar(["Not Affected", "Affected"], counts, color=["green", "red"])
plt.title("CHD Status")
plt.ylabel("Number of People")
plt.xlabel("CHD Status")
plt.show()
```



Splitting the data

```
In [66]: X = np.asarray(disease[['age', 'Sex_male', 'cigsPerDay',
                                'totChol', 'sysBP', 'glucose']])
y = np.asarray(disease['TenYearCHD'])

# normalization of the dataset
X = preprocessing.StandardScaler().fit_transform(X)

# Train test split
from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test = train_test_split(X,y, test_size=0.3, random

print(f"Train set: {X_train.shape} {y_train.shape}")
print(f"Test set: {X_test.shape} {y_test.shape}")
```

Train set: (2968, 6) (2968,)
Test set: (1272, 6) (1272,)

```
In [67]: # As you can see from the barplot there is a bias in data so we use smote to
# Handle class imbalance with SMOTE (Oversampling the minority class)
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X_train, y_train)

print(f"Train set: {X_resampled.shape} {y_resampled.shape}")
print(f"Test set: {X_test.shape} {y_test.shape}")
```

Train set: (5024, 6) (5024,)
Test set: (1272, 6) (1272,)

Fitting Logistic Regression Model for Heart Disease Prediction

```
In [111]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_resampled,y_resampled)
y_pred = model.predict(X_test)
```

Checking the Accuracy of the model

```
In [114]: score_lr=round(accuracy_score(y_test,y_pred)*100,2)
roc_lr= round(roc_auc_score(y_test, model.predict_proba(X_test)[: , 1])*100,
print("Accuracy of the model is =",score_lr)
print("ROC-AUC Score:",roc_lr)
```


Accuracy of the model is = 66.04
ROC-AUC Score: 72.07

```
In [115]: # Confusion matrix
from sklearn.metrics import confusion_matrix, classification_report

cm = confusion_matrix(y_test, y_pred)
conf_matrix = pd.DataFrame(data = cm,
                           columns = ['Predicted:0', 'Predicted:1'],
                           index = ['Actual:0', 'Actual:1'])

plt.figure(figsize = (8, 5))
sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap = "Greens")

plt.show()
print('The details for confusion matrix is =')
print(classification_report(y_test, y_pred))
```



The details for confusion matrix is =

	precision	recall	f1-score	support
0	0.93	0.65	0.77	1084
1	0.26	0.70	0.38	188
accuracy			0.66	1272
macro avg	0.59	0.68	0.57	1272
weighted avg	0.83	0.66	0.71	1272

Fitting Random Forest Classifier for Heart Disease Prediction

```
In [116]: # Train a model
model_RFC = RandomForestClassifier(class_weight='balanced', random_state=42)
model_RFC.fit(X_resampled, y_resampled)

# Evaluate
y_pred = model_RFC.predict(X_test)

print(classification_report(y_test, y_pred))

score_rf = round(accuracy_score(y_test,y_pred)*100,2)
roc_rf = round( roc_auc_score(y_test, model_RFC.predict_proba(X_test)[:,-1])

print("Accuracy of the model: ",score_rf)
print("ROC-AUC Score:",roc_rf)
```

	precision	recall	f1-score	support
0	0.87	0.87	0.87	1084
1	0.27	0.28	0.28	188
accuracy			0.78	1272
macro avg	0.57	0.57	0.57	1272
weighted avg	0.79	0.78	0.79	1272

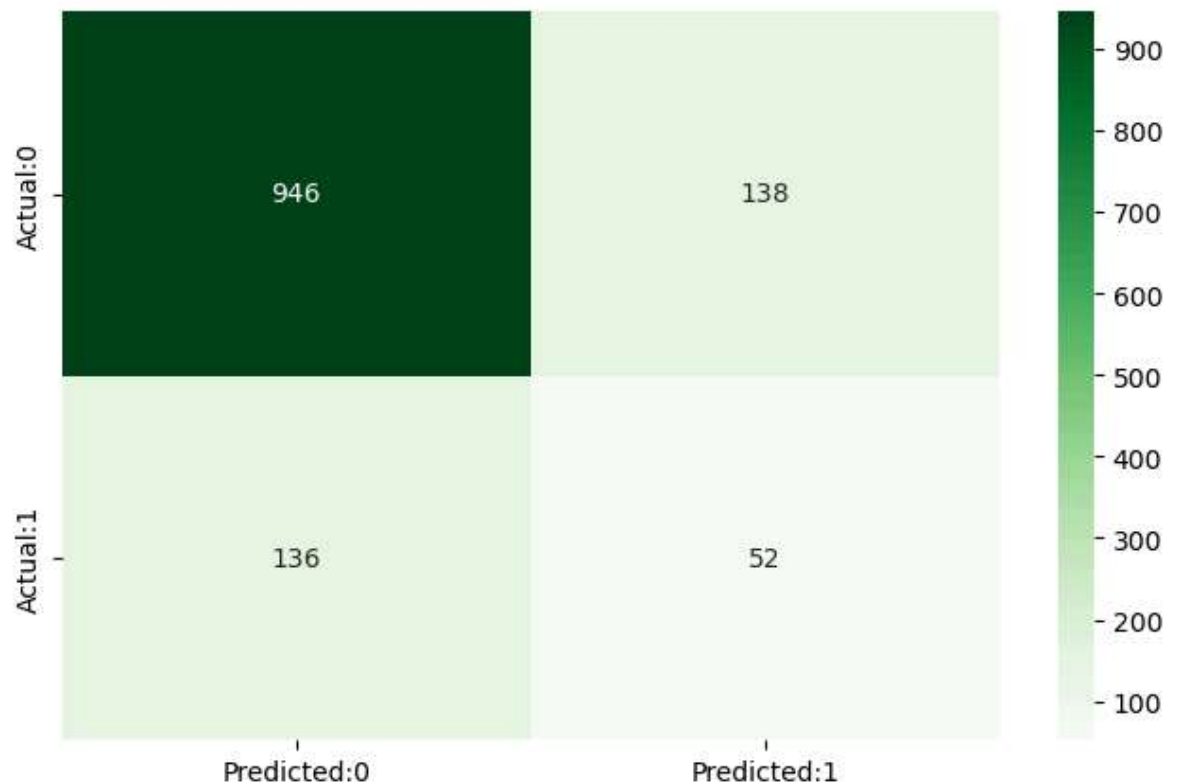
Accuracy of the model: 78.46
 ROC-AUC Score: 68.62

```
In [117]: # Confusion matrix
from sklearn.metrics import confusion_matrix, classification_report

cm = confusion_matrix(y_test, y_pred)
conf_matrix = pd.DataFrame(data = cm,
                           columns = ['Predicted:0', 'Predicted:1'],
                           index = ['Actual:0', 'Actual:1'])

plt.figure(figsize = (8, 5))
sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap = "Greens")

plt.show()
print('The details for confusion matrix is =')
print(classification_report(y_test, y_pred))
```



The details for confusion matrix is =

	precision	recall	f1-score	support
0	0.87	0.87	0.87	1084
1	0.27	0.28	0.28	188
accuracy			0.78	1272
macro avg	0.57	0.57	0.57	1272
weighted avg	0.79	0.78	0.79	1272

Fitting SVC while using gaussian rbf kernel for Heart Disease Prediction

```
In [121]: # Train the SVC model with class weights
svc = SVC(kernel='rbf', class_weight='balanced', probability=True, random_s
svc.fit(X_resampled, y_resampled)

# Make predictions and evaluate the model
y_pred = svc.predict(X_test)

# Evaluate performance
roc_svc_rbf = round(roc_auc_score(y_test, svc.predict_proba(X_test)[: , 1])*
score_svc_rbf = round(accuracy_score(y_test,y_pred)*100,2)

print(classification_report(y_test, y_pred))
print("Acurracy of the model is :",score_svc_rbf)
print("ROC-AUC Score:",roc_svc_rbf)
```

	precision	recall	f1-score	support
0	0.91	0.68	0.78	1084
1	0.25	0.61	0.35	188
accuracy			0.67	1272
macro avg	0.58	0.65	0.57	1272
weighted avg	0.81	0.67	0.72	1272

Acurracy of the model is : 67.3


ROC-AUC Score: 69.62

```
In [122]: # Confusion matrix
from sklearn.metrics import confusion_matrix, classification_report

cm = confusion_matrix(y_test, y_pred)
conf_matrix = pd.DataFrame(data = cm,
                           columns = ['Predicted:0', 'Predicted:1'],
                           index = ['Actual:0', 'Actual:1'])

plt.figure(figsize = (8, 5))
sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap = "Greens")

plt.show()
print('The details for confusion matrix is =')
print(classification_report(y_test, y_pred))
```



The details for confusion matrix is =

	precision	recall	f1-score	support
0	0.91	0.68	0.78	1084
1	0.25	0.61	0.35	188
accuracy			0.67	1272
macro avg	0.58	0.65	0.57	1272
weighted avg	0.81	0.67	0.72	1272

Fitting SVC while using linear kernel for Heart Disease Prediction

```
In [123]: # Train the model
linear_svc_model=SVC(kernel="linear", class_weight="balanced", probability=True)
linear_svc_model.fit(X_resampled,y_resampled)

# Make predictions and evaluate the model
y_pred = linear_svc_model.predict(X_test)

# Evaluate the model
score_linearsvc = round(accuracy_score(y_test,y_pred)*100,2)
roc_linearsvc = round(roc_auc_score(y_test, linear_svc_model.predict_proba(y_test),
                                class_weight="balanced"))
print(classification_report(y_test, y_pred))
print("Accuracy of the model is :",score_linearsvc)
print("ROC-AUC Score:", roc_linearsvc)
```

	precision	recall	f1-score	support
0	0.92	0.65	0.76	1084
1	0.25	0.68	0.36	188
accuracy			0.65	1272
macro avg	0.58	0.66	0.56	1272
weighted avg	0.82	0.65	0.70	1272

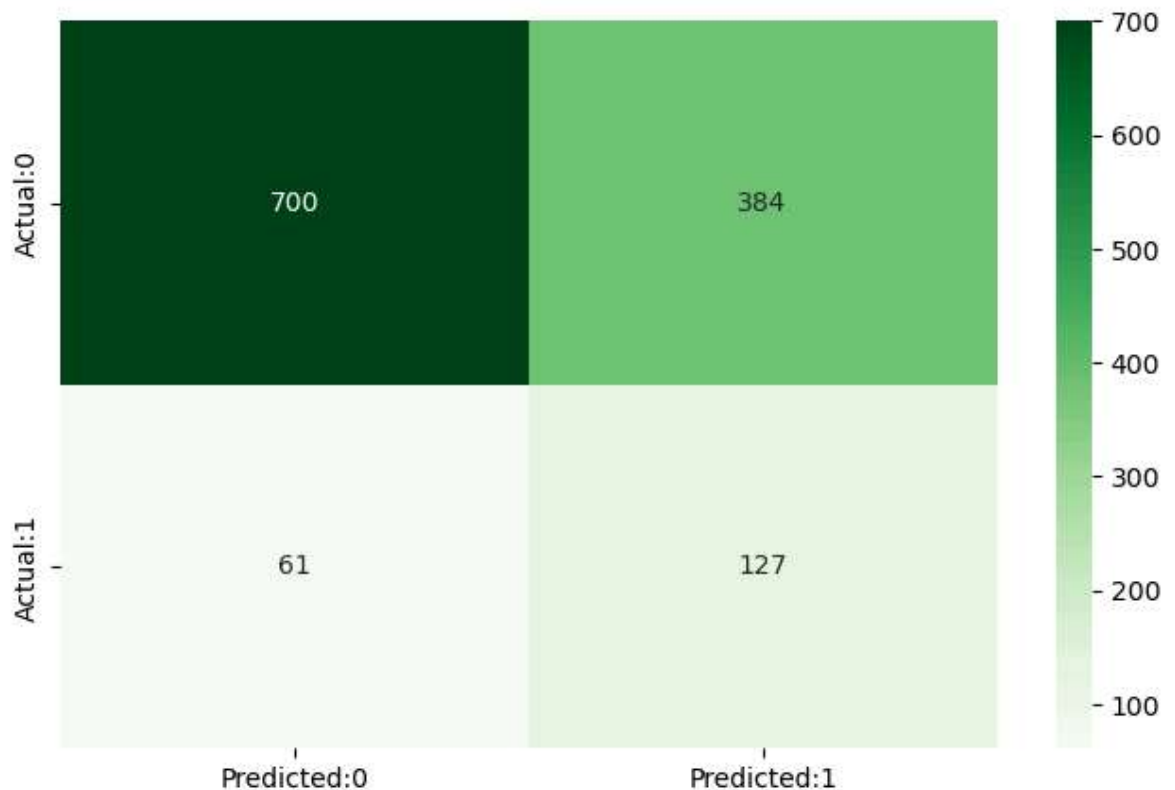
Accuracy of the model is : 65.02
 ROC-AUC Score: 71.87

```
In [124]: # Confusion matrix
from sklearn.metrics import confusion_matrix, classification_report

cm = confusion_matrix(y_test, y_pred)
conf_matrix = pd.DataFrame(data = cm,
                           columns = ['Predicted:0', 'Predicted:1'],
                           index = ['Actual:0', 'Actual:1'])

plt.figure(figsize = (8, 5))
sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap = "Greens")

plt.show()
print('The details for confusion matrix is =')
print(classification_report(y_test, y_pred))
```



The details for confusion matrix is =

	precision	recall	f1-score	support
0	0.92	0.65	0.76	1084
1	0.25	0.68	0.36	188
accuracy			0.65	1272
macro avg	0.58	0.66	0.56	1272
weighted avg	0.82	0.65	0.70	1272

```
In [131]: # Train the model
poly_svc_model=SVC(kernel="poly", degree=4, coef0=1, class_weight="balanced")
poly_svc_model.fit(X_resampled,y_resampled)

# Make predictions and evaluate the model
y_train_predict =poly_svc_model.predict(X_resampled)
y_pred = poly_svc_model.predict(X_test)

# Evaluate the model
score_polysvc = round(accuracy_score(y_test,y_pred)*100,2)
roc_polysvc = round(roc_auc_score(y_test,poly_svc_model.predict_proba(X_test)[:,1]),2)

print(classification_report(y_test, y_pred))
print("Accuracy of the model is :",score_polysvc)
print("ROC-AUC Score:", roc_polysvc)
print("Accuracy of the model on training is :",accuracy_score(y_resampled,y_train_predict))
```

	precision	recall	f1-score	support
0	0.90	0.69	0.78	1084
1	0.23	0.54	0.32	188
accuracy			0.67	1272
macro avg	0.56	0.61	0.55	1272
weighted avg	0.80	0.67	0.71	1272

Accuracy of the model is : 66.9

ROC-AUC Score: 65.18

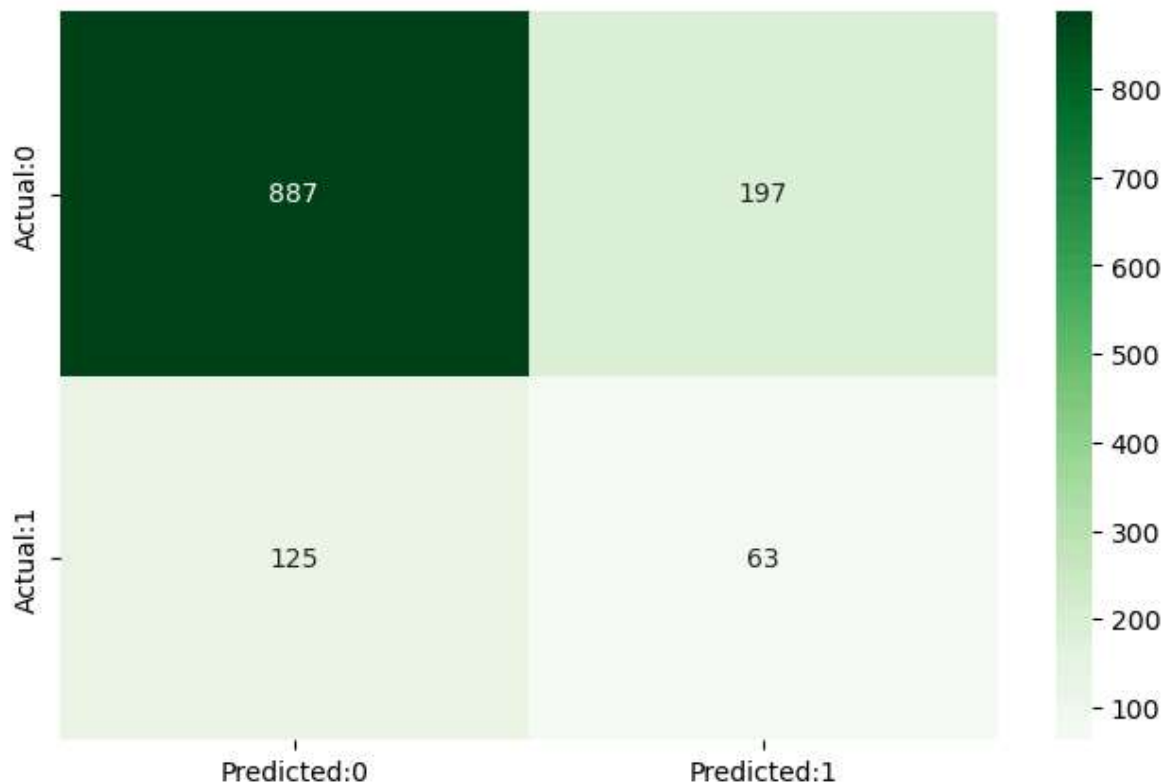
Accuracy of the model on training is : 0.7231289808917197

```
In [109]: # Confusion matrix
from sklearn.metrics import confusion_matrix, classification_report

cm = confusion_matrix(y_test, y_pred)
conf_matrix = pd.DataFrame(data = cm,
                           columns = ['Predicted:0', 'Predicted:1'],
                           index = ['Actual:0', 'Actual:1'])

plt.figure(figsize = (8, 5))
sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap = "Greens")

plt.show()
print('The details for confusion matrix is =')
print(classification_report(y_test, y_pred))
```



The details for confusion matrix is =

	precision	recall	f1-score	support
0	0.88	0.82	0.85	1084
1	0.24	0.34	0.28	188
accuracy			0.75	1272
macro avg	0.56	0.58	0.56	1272
weighted avg	0.78	0.75	0.76	1272

Final Output

```
In [142]: scores= [score_lr,score_linearsvc,score_polysvc,score_rf,score_svc_rbf]
          rocs = [roc_lr,roc_linearsvc,roc_polysvc,roc_rf,roc_svc_rbf]
          algorithms = ["Logistic Regression","LinearSVC","Polynomial SVC","Random Forest","Gaussian RBF SVC"]

          for i in range(len(algorithms)):
              print(f"The Accuracy and ROC score achieved using {algorithms[i]} is: {scores[i]} and {rocs[i]}")
```

The Accuracy and ROC score achieved using Logistic Regression is: 66.04% and 72.07%

The Accuracy and ROC score achieved using LinearSVC is: 65.02% and 71.87%


The Accuracy and ROC score achieved using Polynomial SVC is: 66.9% and 65.18%

The Accuracy and ROC score achieved using Random Forest is: 78.46% and 68.62%

The Accuracy and ROC score achieved using Gaussian RBF SVC is: 67.3% and 69.62%

```
In [145]: sns.set(rc={'figure.figsize': (15, 8)})
          plt.xlabel("Algorithms")
          plt.ylabel("Accuracy Score")

          # Explicitly specify x and y
          sns.barplot(x=algorithms, y=scores)
          plt.show()
```



In []:

